

Understanding Image Advertisements and Predicting Sentiment

1st Rohan Chopra
40233019
Concordia University
Montreal, Canada
rohanchopra09@gmail.com

2nd Harman Singh Jolly
40204947
Concordia University
Montreal, Canada
jollyharmansingh@gmail.com

Abstract—This document is a model and instructions for L^AT_EX. This and the IEEETran.cls file define the components of your paper [title, text, heads, etc.]. *CRITICAL: Do Not Use Symbols, Special Characters, Footnotes, or Math in Paper Title or Abstract.

Index Terms—component, formatting, style, styling, insert

I. Introduction

In this fast paced world, companies now trade with the ever changing data we generate. With data, comes along user profiling or a set of topics that a user can be defined into. Perhaps, the expansion of online visual content and social media has led to a surge of interest in the investigation of large-scale social multimedia analysis. In recent years, image advertisements have become increasingly prevalent in our daily lives. From billboards to social media ads, companies are constantly vying for our attention through visually appealing images. However, with the rise of sentiment analysis, it is now possible to predict the emotions that these images evoke in viewers. So, pondering over the platitude, "A picture says a thousand words" becomes necessary.

Tech giants like Facebook and Google have produced enormous wealth by advertisements only where 97.5% of the \$116.6 billion generated by Facebook [1] and 70.9% of Google's revenue [2] came in from broadcasting advertisements. Attracting customers to a product is a vital aspect of using advertisements on both television and social media. Selection of ads by hand is a physically tedious and onerous job where provision of image based advertisements or videos becomes highly difficult via handpicking. So, automatic advertising techniques are developed, such as the contextual advertising method which aims to find the most relevant ad to a provided content without annoying customers.

The components that contribute to the effectiveness of these advertisements are intricate and are being analyzed in marketing science and consumer psychology. A particular focus of this research has been on comprehending the emotions and sentiments conveyed by visual media content, which has become increasingly popular for both academic and practical purposes. Sentiment analysis is a process of identifying and extracting opinions, emotions,

and attitudes expressed in text. However, with the increasing use of visual media in advertising, there has been a need for sentiment analysis in images. Sentiment labels are used to annotate images with positive, negative or neutral sentiments. These annotations are used to train machine learning models to predict the emotions that an image is likely to evoke in viewers. Powerful emotions conveyed by images and videos can amplify the message conveyed in the content, making it more impactful and capable of influencing the audience more effectively. But sometimes, advertisements aren't able to hold viewers interest as the product shown doesn't necessarily incline to the viewers choices or viewers actually skip watching advertisements that aren't able to garner their attention, which is where a positive or a negative connotation comes into picture pertaining to sentiment of an advertisement.

Our report focuses on .

II. Related Works

In order to create an effective advertisement, researchers focus on tons of ideas to reach a balance of emotion and the type of message conveyed. The visual understanding or common sense that the "red color" symbolizes "stop" when used in traffic related series and symbolizes "blood" when used in medical related series is extremely potent in humans but negligible in computers. The recognition of non-photorealistic objects and symbolism are possible illustrations depicting the visual-rhetoric methods for conveying a message through ads. To fathom these rhetorics which not only requires the discernment of objects but also requires the decoding of this rhetoric [4] [5]. Perhaps, a summarization about the image can where it talks about [6] [7] which objects are portrayed, how they are portrayed and answering the primal questions as to why they are portrayed. For instance, if an advertisement is meant to target a young audience, it can be designed with bright colors and positive sentiments to appeal to that demographic. On the other hand, an ad meant for an older audience may be designed with muted colors and neutral sentiments. By predicting the sentiments that an advertisement is likely to evoke, companies can tailor

their ads to specific audiences, increasing their chances of success.

Exploring objects or generic nouns such as “cat” or “trees” have been a marvel of machine vision but the association of sentiments correlating to the visuals remains a challenging and seemingly insurmountable task. The difficulty of this endeavor stems from the significant emotional distance between the low visual characteristics and the high-level sentiment that adjectives convey. To fill this gap between visual characteristics and sentiments, Borth, Damian et al. [8] proposed an alternative method that represents sentiment-related visual concepts for an intermediary representation. They have coined the term Adjective Noun Pairs (ANP) like “sprawling trees” or “sleepy cat” which are used to corroborate the adjective of nouns, further used to identify nouns. Although ANPs do not directly express emotions or sentiments, they serve as useful statistical cues for detecting emotions conveyed in the images, as they were discovered based on strong co-occurrence relationships with emotion tags of web photos. In their study Borth, Damian et al [8] trained binary SVM classifiers on ANPs for whole images. To further enhance the classifiers, Chen, Felix et al. [9] incorporated object-based concept localization and the semantic similarity among the concepts.

DeepSentiBank [10] shows significant efforts to interpret emotion from images using a model trained on web images that are tagged with multiple sentiments. Hussain, Zhang et al. [11] went further to apply DeepSentiBank and reached a conclusion to observe a detector for natural images (DeepSentiBank) couldn't be applicable for advertisement images. Besides DeepSentiBank, Vedula & Sun et al. [12] developed an advertisement recommendation system using sentiments in multimedia content. Using deep learning techniques, image and video analysis combined have garnered a great amount of headway for data analysis at YouTube and for contextual understanding in video advertising [13]. But, compared to other forms of media such as publishing or billboards, which are less personalized, a great deal of lacuna lies in interpreting advertisements media due to their extensive use of visual rhetoric as coined in Hussain, Zhang et al. [11].

An advertisement demands much in depth analysis due to the fact that they intend to persuade people for their motion by lobbying their product as best to make people aware of any social cause. Give an example of ad images here, some advertisement that represents a hidden social cause. The motive to persuade people is deeply imbibed inside an advertisement, where sometimes it is portrayed via simple tone or sometimes via sarcasm. In order to understand this underlying sentiment, only appearance isn't enough. Therefore, to personalize the experience with ads, it is imperative to not only understand its topic, but also the emotion it conveys. Zhang, Luo et al. [14] introduced a framework that amalgamates multiple modalities together for envisioning topic and sentiment related predictions to

further conceptualize advertisements.

III. Dataset

Before you begin to format your paper, first write and save the content as a separate text file. Complete all content and organizational editing before formatting. Please note sections ?? below for more information on proofreading, spelling and grammar.

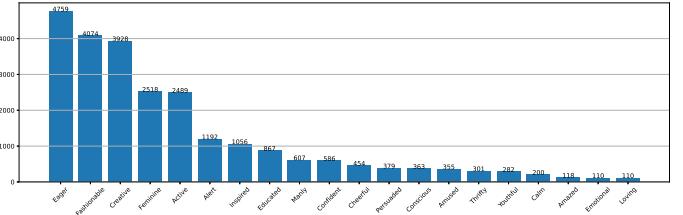


Fig. 1. Sentiment Label Distribution

IV. Methodology

Before training, the images have to be processed for denoising and their quality improved before loading it to model for classification.

A. Bilateral Filter

For improving the quality of the images, a bilateral filter was used. A smoothing filter for images, is used to reduce noise while preserving edges in a non-linear manner. This filter works by calculating a weighted average of intensity values from surrounding pixels to replace the intensity of each pixel. The weight of each pixel is determined using a Gaussian distribution and it preserves sharp edges. The weights used in the filter are not solely based on the distance between pixels, but also take into consideration differences in radiometric properties such as color intensity and depth distance [15].

As shown in fig 2, a bilateral filter is controlled by three parameters sigma(s), responsible for spatial region which smooths larger surfaces as we increase the spatial parameter, sigma(r) which tends to act like a Gaussian filter when it increases as the Gaussian range widens and d, which is the diameter of each pixel neighbourhood. It is quintessential to know that all other filters smudge the edges, while Bilateral Filtering retains them. When weights are multiplied together and if one of the weights is close to zero, no smoothing occurs. This can be demonstrated by using a large spatial Gaussian with a narrow range Gaussian, which results in limited smoothing despite the filter having a large spatial extent. The range weight ensures that the edges are preserved.

After applying different values d, sigma(s) and sigma(r), it was determined that the pixel diameter of 9 along with sigma(s) and sigma(r) being 9 as well gave the best results. Following that, bilateral filters have been applied twice merely because of the fact that applying bilateral filters in iterations enhanced picture quality even more. The

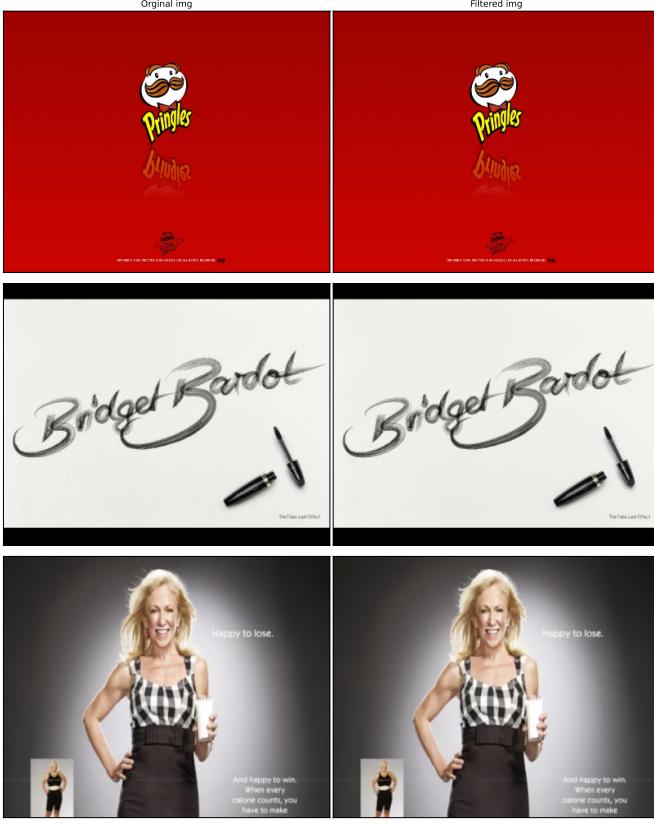


Fig. 2. Bilateral Filter applied on Ads

image below depicts a significant change than the previous image i.e. the original image where the color banding or posterization, an ugly artifact that can be seen in digital images, has been notably reduced [16] and the final image is better than previous one. Another remarkable change was also witnessed which relates to compression artifacts, a distortion of media in images which is caused by lossy compression of media was removed in the final image corroborating the efforts.

B. Pre-processing Techniques

Pytorch provides various functional transformations that can be applied using the `torchvision.transforms` module. They accept both PIL images and tensor images, although some transformations are PIL-only and some are tensor-only [?]. As these transformations require a parameter such as a factor by which an image can be transformed, therefore they cannot be applied to all images owing to the factor that all images are different.

For example, a Hue transform that accepts an image along with a parameter, `hue_factor` that ranges from [-0.5 to 0.5]. 0.5 and -0.5 give complete reversal of the hue channel in HSV space in positive and negative direction respectively whereas 0 means no shift. The same level factor cannot be expected from other techniques such as sharpness or contrast etc. Hence, this parameter cannot

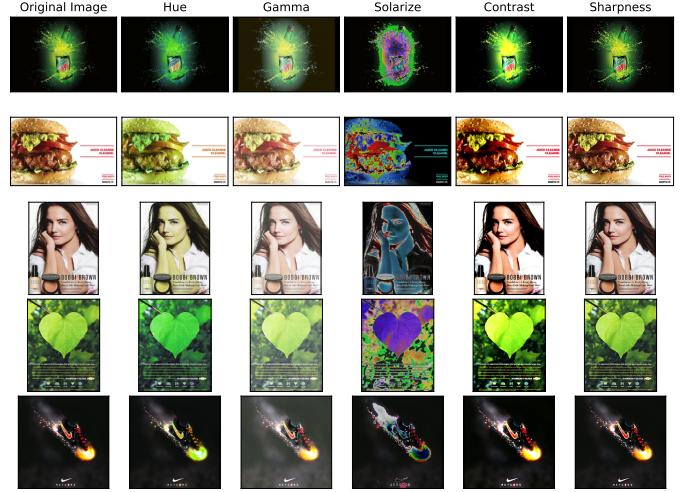


Fig. 3. Pytorch Preprocessing Techniques

be kept constant for all images as it will have a variable effect or appeal on different images. In the fig below, five random images have been selected and functional image processing techniques like hue transforms, gamma transforms, solarize transformations, sharpness etc have been applied to reach a conclusion that all images bearing unique in their characteristics respond differently to functional transformations applied.

Even though brightness and contrast change show a potent outcome, a specific parameter cannot be kept for all images which points this approach to be handled in data augmentation where functions like AutoContrast and AutoBrightness have been applied. Histogram equalization, a popular technique for improving the contrast of the images visibly failed to work for colored images. This technique was employed on a colored image which visibly distorted colors and features such that it could not be considered as a viable pre-processing technique.



Fig. 4. Histogram Equalization on Colored Image

Some details about an image comparison between image 1 and image 2....such that we are able to substantiate our

findings.

C. Data Augmentations Techniques

Data augmentation plays a crucial role in completing the training requirements for deep learning models as these models aren't able to converge the network to an optimal solution using some gradient based optimization algorithm. The huge number of parameters needed to be tuned by the learning algorithm require enormous amounts of data merely because of the fact that the deep learning algorithms start off with a poor initial state where weights are completely random. There are various ways to augment data using pytorch library such as RandomHorizontalFlip, RandomAdjustSharpness etc that produce images with any random factor for creating possible image outputs that may be a possible scenario in the real world for understanding the variability in the image dataset. In the images below, data augmentation has been visualized using RandomHorizontalFlip, RandomRotation to a certain degree, RandomAdjustSharpness with sharpness factor as 2, RandomAutocontrast with factor as 5, and lastly RandomColorJitter.



Fig. 5. Data Augmentation using RandomHorizontalFlip



Fig. 6. Data Augmentation using RandomRotation at 30 degrees



Fig. 7. Data Augmentation using ColorJitter

D. Model Architectures

(I) Different backbone architectures were chosen to ensure that different types of Convolution blocks were tested for the advertisement data. Other selection criteria



Fig. 8. Data Augmentation using RandomAutocontrast



Fig. 9. Data Augmentation using RandomAdjustSharpness

included the number of parameters and GFLOPS, important to keep track of the total training and evaluation time, and the top 5 classification accuracy on the ImageNet 1K benchmark dataset. Finally, the following three backbone architectures were chosen:

TABLE I
Shortlisted Backbone Architectures.

Arch.	Params (Mil.)	Layers	GFLOPS	Imagenet Acc.
MobileNet V3 Large	5.5	18	0.22	92.57
EfficientNet B3	12.2	29	1.83	96.05
Resnet 50	25.6	50	4.09	95.43

ResNet 50: A residual learning CNN with 50 layers that are made possible by skip connections. Without these skip connections, training such a deep network is not possible due to the vanishing gradient problem. The 50 layer variant was chosen to decrease training time while not compromising on the accuracy. This architecture had the highest trainable parameters, FLOPS and number of layers [18].

MobileNet V3 Large: This model uses depthwise separable convolution from MobileNet V2 along with squeeze-excitation blocks in residual layers from MnasNet. This makes it really quick to train while still performing at par with other architectures. This architecture had the lowest trainable parameters and FLOPS among the three selected. Howard et al. [19] also used network architecture search to find the most effective model. The large configuration was chosen to not compromise on the prediction accuracy.

EfficientNet B3: This model uses compound scaling to scale the model by depth, width and resolution. The B3 version was chosen to have faster training without compromising on the accuracy. [20] This architecture performs the best among the selected on the Imagenet benchmark

dataset while having half the trainable parameters of Resnet50.

E. Optimization Algorithm

The Adam optimizer [?] is an adaptive learning rate optimization algorithm which was chosen as the optimizer for this study as it converges faster by integrating benefits of the RMSProp algorithm and momentum technique. It is also robust to hyperparameters but, requires tweaking of the learning rate depending on the task at hand. For this study, a learning rate of 0.01 and the original author recommend settings for $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-8}$ were used for the first and second order moment estimate as defined in (1) and (2) where β_1 and β_2 control the decay rates.

$$m_t = \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t \quad (1)$$

$$v_t = \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2 \quad (2)$$

Further, the Cosine annealing [?] learning rate scheduler was used to reduce the learning rate as the training progressed down to an end point of 0.001. This would help us reduce the learning rate as training progressed, preventing us from overshooting the minima.

V. Results

A. Experiment Setup

First, training data with labels that were low in count were removed as described in ???. Then, the advertisement images were pre-processed using the bilateral filter described in ?? before resizing them to the size - 384 * 384. The data was split into train, test and validation set in the 0.6:0.2:0.2 ratio and was stored in separate directories according to the defined PyTorch dataloader. Then, the mean and standard deviation of the dataset was calculated using the training dataset. All images were normalized before training using the dataloader with this calculated mean and standard deviation. The dataset used in this study presented the the multiclass, multilabel classification problem. Thus, to make the model predict multiple labels, a sigmoid layer had to be added before the loss function to get 0 or 1 prediction for all the classes of the data. To achieve this, the BCEWITHLOGITSLOSS function of PyTorch was used as it combines the Sigmoid layer and the binary cross entropy loss function in one single class. This makes theses operations more numerically stable than their separate counterparts [?]. The backbone architectures and their pre-trained weights were obtained directly from the torchvision library and the final classification layer was modified for our dataset of 20 classes. The pre-trained weights were chosen to be the IMAGENET1K_V2 weights and only the last classification layer was fine-tuned. The rationale behind performing this type of shallow-tuning was that the Imagenet data is very similar to the advertisement images in our dataset. Additionally, the size of the selected dataset is small so deep-tuning might not work well.

The batch size was fixed to 32 for all the models. While training, the best model by validation loss was saved to prevent the usage of overfit models for the test set analysis. The actual and predicted results from each epoch was also stored to calculate the F1 scores at each step of training. While calculating the F1 score, macro averaging was used to get an average score across classes.

Initial training runs of the multilabel data produced a zero F1 score due to its highly imbalanced nature. To mitigate this, class wise weights were calculated and used with the loss function. This improved the F1 score quite considerably.

TABLE II

F1 (higher is better), time per epoch in seconds (lower is better), and number of epochs to reach the best validation loss (lower is better) for the 3 models that were trained.

Model	F1	Time	Epochs
MobileNet V3 Large	0.168	80s	abc
EfficientNet B3	0.189	153s	abc
Resnet 50	0.179	89s	abc

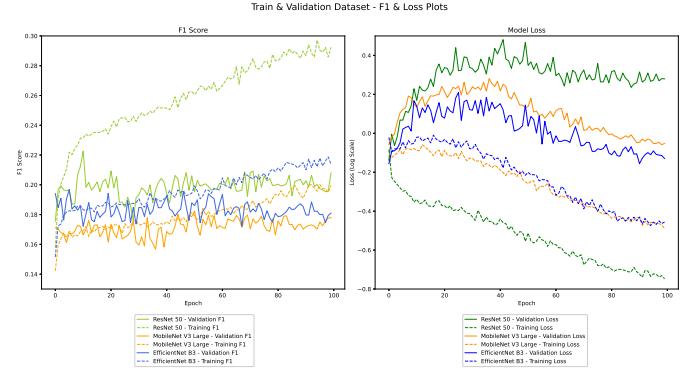


Fig. 10. Training & Validation, F1 & Loss plots for the three models.

Finally, the best models from each of the three training runs validation loss were used to get the test set metrics that are displayed in ???. Per epoch training and validation F1 score and loss are provided in Fig. 10.

B. Training Results

References

- [1] <https://www.oberlo.com/statistics/facebook-ad-revenue>
- [2] <https://www.statista.com/statistics/266249/advertising-revenue-of-google/>
- [3] Mei, Tao, Xian-Sheng Hua, Linjun Yang, and Shipeng Li. "VideoSense: towards effective online video advertising." In Proceedings of the 15th ACM international conference on Multimedia, pp. 1075-1084. 2007.
- [4] Girshick, Ross. "Fast r-cnn." In Proceedings of the IEEE international conference on computer vision, pp. 1440-1448. 2015.
- [5] Szegedy, Christian, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott Reed, Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, and Andrew Rabinovich. "Going deeper with convolutions." In Proceedings of the IEEE conference on computer vision and pattern recognition, pp. 1-9. 2015.

- [6] Y. Donahue, Jeffrey, Lisa Anne Hendricks, Sergio Guadarrama, Marcus Rohrbach, Subhashini Venugopalan, Kate Saenko, and Trevor Darrell. "Long-term recurrent convolutional networks for visual recognition and description." In Proceedings of the IEEE conference on computer vision and pattern recognition, pp. 2625-2634. 2015.
- [7] Vinyals, Oriol, Alexander Toshev, Samy Bengio, and Dumitru Erhan. "Show and tell: A neural image caption generator." In Proceedings of the IEEE conference on computer vision and pattern recognition, pp. 3156-3164. 2015.
- [8] Borth, Damian, Rongrong Ji, Tao Chen, Thomas Breuel, and Shih-Fu Chang. "Large-scale visual sentiment ontology and detectors using adjective noun pairs." In Proceedings of the 21st ACM international conference on Multimedia, pp. 223-232. 2013.
- [9] Chen, Tao, Felix X. Yu, Jiawei Chen, Yin Cui, Yan-Ying Chen, and Shih-Fu Chang. "Object-based visual sentiment concept analysis and application." In Proceedings of the 22nd ACM international conference on Multimedia, pp. 367-376. 2014.
- [10] Chen, Tao, Damian Borth, Trevor Darrell, and Shih-Fu Chang. "Deepsentibank: Visual sentiment concept classification with deep convolutional neural networks." arXiv preprint arXiv:1410.8586 (2014).
- [11] Hussain, Zaeem, Mingda Zhang, Xiaozhong Zhang, Keren Ye, Christopher Thomas, Zuhra Agha, Nathan Ong, and Adriana Kovashka. "Automatic understanding of image and video advertisements." In Proceedings of the IEEE conference on computer vision and pattern recognition, pp. 1705-1715. 2017.
- [12] Vedula, Nikhita, Wei Sun, Hyunhwan Lee, Harsh Gupta, Mitsunori Ogihara, Joseph Johnson, Gang Ren, and Srinivasan Parthasarathy. "Multimodal content analysis for effective advertisements on youtube." In 2017 IEEE international conference on data mining (ICDM), pp. 1123-1128. IEEE, 2017.
- [13] Madhok, Rishi, Shashank Majumdar, Nitin Gupta, and Sameep Mehta. "Semantic understanding for contextual in-video advertising." In Proceedings of the AAAI Conference on Artificial Intelligence, vol. 32, no. 1. 2018.
- [14] Zhang, Huaizheng, Yong Luo, Qiming Ai, Yonggang Wen, and Han Hu. "Look, read and feel: Benchmarking ads understanding with multimodal multitask learning." In Proceedings of the 28th ACM International Conference on Multimedia, pp. 430-438. 2020.
- [15] https://en.wikipedia.org/wiki/Bilateral_filter
- [16] <https://www.willgibbons.com/color-banding/>
- [17] <https://pytorch.org/vision/stable/transforms.html>
- [18] He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. "Deep residual learning for image recognition." In Proceedings of the IEEE conference on computer vision and pattern recognition, pp. 770-778. 2016.
- [19] Howard, Andrew, Mark Sandler, Grace Chu, Liang-Chieh Chen, Bo Chen, Mingxing Tan, Weijun Wang et al. "Searching for mobilenetv3." In Proceedings of the IEEE/CVF international conference on computer vision, pp. 1314-1324. 2019.
- [20] Tan, Mingxing, and Quoc Le. "Efficientnet: Rethinking model scaling for convolutional neural networks." In International conference on machine learning, pp. 6105-6114. PMLR, 2019.